
schedule Documentation

Release 0.4.0

Daniel Bader

Jan 23, 2019

Contents

1	Features	3
2	Usage	5
3	API Documentation	7
3.1	Developer Interface	7
4	Common Questions	11
4.1	Frequently Asked Questions	11
5	Issues	17
6	About Schedule	19
	Python Module Index	21

Python job scheduling for humans.

An in-process scheduler for periodic jobs that uses the builder pattern for configuration. Schedule lets you run Python functions (or any other callable) periodically at pre-determined intervals using a simple, human-friendly syntax.

Inspired by [Adam Wiggins](#)' article "[Rethinking Cron](#)" and the [clockwork](#) Ruby module.

CHAPTER 1

Features

- A simple to use API for scheduling jobs.
- Very lightweight and no external dependencies.
- Excellent test coverage.
- Tested on Python 2.7, 3.5, and 3.6

CHAPTER 2

Usage

```
$ pip install schedule
```

```
import schedule
import time

def job():
    print("I'm working...")

schedule.every(10).minutes.do(job)
schedule.every().hour.do(job)
schedule.every().day.at("10:30").do(job)
schedule.every().monday.do(job)
schedule.every().wednesday.at("13:15").do(job)
schedule.every().minute.at(":17").do(job)

while True:
    schedule.run_pending()
    time.sleep(1)
```


If you are looking for information on a specific function, class, or method, this part of the documentation is for you.

3.1 Developer Interface

This part of the documentation covers all the interfaces of `schedule`. For parts where `schedule` depends on external libraries, we document the most important right here and provide links to the canonical documentation.

3.1.1 Main Interface

`schedule.default_scheduler = <schedule.Scheduler object>`
Default *Scheduler* object

`schedule.jobs = []`
Default *Jobs* list

`schedule.every(interval=1)`
Calls *every* on the *default scheduler instance*.

`schedule.run_pending()`
Calls *run_pending* on the *default scheduler instance*.

`schedule.run_all(delay_seconds=0)`
Calls *run_all* on the *default scheduler instance*.

`schedule.clear(tag=None)`
Calls *clear* on the *default scheduler instance*.

`schedule.cancel_job(job)`
Calls *cancel_job* on the *default scheduler instance*.

`schedule.next_run()`
Calls *next_run* on the *default scheduler instance*.

`schedule.idle_seconds()`

Calls `idle_seconds` on the *default scheduler instance*.

3.1.2 Exceptions

exception `schedule.CancelJob`

Can be returned from a job to unschedule itself.

3.1.3 Classes

class `schedule.Scheduler`

Objects instantiated by the *Scheduler* are factories to create jobs, keep record of scheduled jobs and handle their execution.

run_pending()

Run all jobs that are scheduled to run.

Please note that it is *intended behavior that run_pending() does not run missed jobs*. For example, if you've registered a job that should run every minute and you only call `run_pending()` in one hour increments then your job won't be run 60 times in between but only once.

run_all (*delay_seconds=0*)

Run all jobs regardless if they are scheduled to run or not.

A delay of *delay* seconds is added between each job. This helps distribute system load generated by the jobs more evenly over time.

Parameters *delay_seconds* – A delay added between every executed job

clear (*tag=None*)

Deletes scheduled jobs marked with the given tag, or all jobs if tag is omitted.

Parameters *tag* – An identifier used to identify a subset of jobs to delete

cancel_job (*job*)

Delete a scheduled job.

Parameters *job* – The job to be unscheduled

every (*interval=1*)

Schedule a new periodic job.

Parameters *interval* – A quantity of a certain time unit

Returns An unconfigured *Job*

next_run

Datetime when the next job should run.

Returns A datetime object

idle_seconds

Returns Number of seconds until *next_run*.

class `schedule.Job` (*interval, scheduler=None*)

A periodic job as used by *Scheduler*.

Parameters

- **interval** – A quantity of a certain time unit

- **scheduler** – The *Scheduler* instance that this job will register itself with once it has been fully configured in *Job.do()*.

Every job runs at a given fixed time interval that is defined by:

- a *time unit*
- a quantity of *time units* defined by *interval*

A job is usually created and returned by *Scheduler.every()* method, which also defines its *interval*.

second

seconds

minute

minutes

hour

hours

day

days

week

weeks

monday

tuesday

wednesday

thursday

friday

saturday

sunday

tag (*tags)

Tags the job with one or more unique indentifiers.

Tags must be hashable. Duplicate tags are discarded.

Parameters **tags** – A unique list of Hashable tags.

Returns The invoked job instance

at (time_str)

Specify a particular time that the job should be run at.

Parameters **time_str** – A string in one of the following formats: *HH:MM:SS*, *HH:MM*, *‘:MM’*, *:SS*. The format must make sense given how often the job is repeating; for example, a job that repeats every minute should not be given a string in the form *HH:MM:SS*. The difference between *:MM* and *:SS* is inferred from the selected time-unit (e.g. *every().hour.at(‘:30’)* vs. *every().minute.at(‘:30’)*).

Returns The invoked job instance

to (latest)

Schedule the job to run at an irregular (randomized) interval.

The job's interval will randomly vary from the value given to *every* to *latest*. The range defined is inclusive on both ends. For example, *every(A).to(B).seconds* executes the job function every N seconds such that $A \leq N \leq B$.

Parameters **latest** – Maximum interval between randomized job runs

Returns The invoked job instance

do (*job_func*, **args*, ***kwargs*)

Specifies the *job_func* that should be called every time the job runs.

Any additional arguments are passed on to *job_func* when the job runs.

Parameters **job_func** – The function to be scheduled

Returns The invoked job instance

should_run

Returns `True` if the job should be run now.

run ()

Run the job and immediately reschedule it.

Returns The return value returned by the *job_func*

Please check [here](#) before creating a new issue ticket.

4.1 Frequently Asked Questions

Frequently asked questions on the usage of `schedule`.

4.1.1 How to execute jobs in parallel?

I am trying to execute 50 items every 10 seconds, but from the my logs it says it executes every item in 10 second schedule serially, is there a work around?

By default, `schedule` executes all jobs serially. The reasoning behind this is that it would be difficult to find a model for parallel execution that makes everyone happy.

You can work around this restriction by running each of the jobs in its own thread:

```
import threading
import time
import schedule

def job():
    print("I'm running on thread %s" % threading.current_thread())

def run_threaded(job_func):
    job_thread = threading.Thread(target=job_func)
    job_thread.start()

schedule.every(10).seconds.do(run_threaded, job)
```

(continues on next page)

(continued from previous page)

```
schedule.every(10).seconds.do(run_threaded, job)
schedule.every(10).seconds.do(run_threaded, job)
schedule.every(10).seconds.do(run_threaded, job)
schedule.every(10).seconds.do(run_threaded, job)

while 1:
    schedule.run_pending()
    time.sleep(1)
```

If you want tighter control on the number of threads use a shared jobqueue and one or more worker threads:

```
import Queue
import time
import threading
import schedule

def job():
    print("I'm working")

def worker_main():
    while 1:
        job_func = jobqueue.get()
        job_func()
        jobqueue.task_done()

jobqueue = Queue.Queue()

schedule.every(10).seconds.do(jobqueue.put, job)
schedule.every(10).seconds.do(jobqueue.put, job)
schedule.every(10).seconds.do(jobqueue.put, job)
schedule.every(10).seconds.do(jobqueue.put, job)
schedule.every(10).seconds.do(jobqueue.put, job)

worker_thread = threading.Thread(target=worker_main)
worker_thread.start()

while 1:
    schedule.run_pending()
    time.sleep(1)
```

This model also makes sense for a distributed application where the workers are separate processes that receive jobs from a distributed work queue. I like using `beanstalkd` with the `beanstalkc` Python library.

4.1.2 How to continuously run the scheduler without blocking the main thread?

Run the scheduler in a separate thread. Mrwhick wrote up a nice solution in to this problem [here](#) (look for `run_continuously()`)

4.1.3 Does schedule support timezones?

Vanilla `schedule` doesn't support timezones at the moment. If you need this functionality please check out @imiric's work [here](#). He added timezone support to `schedule` using `python-dateutil`.

4.1.4 What if my task throws an exception?

Schedule doesn't catch exceptions that happen during job execution. Therefore any exceptions thrown during job execution will bubble up and interrupt schedule's `run_xyz` function.

If you want to guard against exceptions you can wrap your job function in a decorator like this:

```
import functools

def catch_exceptions(cancel_on_failure=False):
    def catch_exceptions_decorator(job_func):
        @functools.wraps(job_func)
        def wrapper(*args, **kwargs):
            try:
                return job_func(*args, **kwargs)
            except:
                import traceback
                print(traceback.format_exc())
                if cancel_on_failure:
                    return schedule.CancelJob
        return wrapper
    return catch_exceptions_decorator

@catch_exceptions(cancel_on_failure=True)
def bad_task():
    return 1 / 0

schedule.every(5).minutes.do(bad_task)
```

Another option would be to subclass Schedule like @mplewis did in [this example](#).

4.1.5 How can I run a job only once?

```
def job_that_executes_once():
    # Do some work ...
    return schedule.CancelJob

schedule.every().day.at('22:30').do(job_that_executes_once)
```

4.1.6 How can I cancel several jobs at once?

You can cancel the scheduling of a group of jobs selecting them by a unique identifier.

```
def greet(name):
    print('Hello {}'.format(name))

schedule.every().day.do(greet, 'Andrea').tag('daily-tasks', 'friend')
schedule.every().hour.do(greet, 'John').tag('hourly-tasks', 'friend')
schedule.every().hour.do(greet, 'Monica').tag('hourly-tasks', 'customer')
schedule.every().day.do(greet, 'Derek').tag('daily-tasks', 'guest')

schedule.clear('daily-tasks')
```

Will prevent every job tagged as `daily-tasks` from running again.

4.1.7 I'm getting an `AttributeError`: 'module' object has no attribute 'every' when I try to use schedule. How can I fix this?

This happens if your code imports the wrong `schedule` module. Make sure you don't have a `schedule.py` file in your project that overrides the `schedule` module provided by this library.

4.1.8 How can I add generic logging to my scheduled jobs?

The easiest way to add generic logging functionality to your `schedule` job functions is to implement a decorator that handles logging in a reusable way:

```
import functools
import time

import schedule

# This decorator can be applied to
def with_logging(func):
    @functools.wraps(func)
    def wrapper(*args, **kwargs):
        print('LOG: Running job "%s"' % func.__name__)
        result = func(*args, **kwargs)
        print('LOG: Job "%s" completed' % func.__name__)
        return result
    return wrapper

@with_logging
def job():
    print('Hello, World.')

schedule.every(3).seconds.do(job)

while 1:
    schedule.run_pending()
    time.sleep(1)
```

4.1.9 How to run a job at random intervals?

```
def my_job():
    # This job will execute every 5 to 10 seconds.
    print('Foo')

schedule.every(5).to(10).seconds.do(my_job)
```

4.1.10 How can I pass arguments to the job function?

`do()` passes extra arguments to the job function:

```
def greet(name):
    print('Hello', name)
```

(continues on next page)

(continued from previous page)

```
schedule.every(2).seconds.do(greet, name='Alice')
schedule.every(4).seconds.do(greet, name='Bob')
```

4.1.11 How can I make sure long-running jobs are always executed on time?

Schedule does not account for the time it takes the job function to execute. To guarantee a stable execution schedule you need to move long-running jobs off the main-thread (where the scheduler runs). See “How to execute jobs in parallel?” in the FAQ for a sample implementation.

CHAPTER 5

Issues

If you encounter any problems, please [file an issue](#) along with a detailed description. Please also check the [Frequently Asked Questions](#) and use the search feature in the issue tracker beforehand to avoid creating duplicates. Thank you

CHAPTER 6

About Schedule

Schedule was created by Daniel Bader - @dbader_org

Distributed under the MIT license. See `LICENSE.txt` for more information.

Thanks to all the wonderful folks who have contributed to schedule over the years:

- mattss <<https://github.com/mattss>>
- mrhwick <<https://github.com/mrhwick>>
- cfrco <<https://github.com/cfrco>>
- matrixise <<https://github.com/matrixise>>
- abultman <<https://github.com/abultman>>
- mplewis <<https://github.com/mplewis>>
- WoLfulus <<https://github.com/WoLfulus>>
- dylwhich <<https://github.com/dylwhich>>
- fkromer <<https://github.com/fkromer>>
- alaingilbert <<https://github.com/alaingilbert>>
- Zerrossetto <<https://github.com/Zerrossetto>>
- yetingsky <<https://github.com/yetingsky>>
- schnepp <<https://github.com/schnepp>> <<https://bitbucket.org/saschaschnepp>>
- grampajoe <<https://github.com/grampajoe>>
- gilbsgilbs <<https://github.com/gilbsgilbs>>
- Nathan Wailes <<https://github.com/NathanWailes>>
- Connor Skees <<https://github.com/ConnorSkees>>

S

`schedule`, [7](#)

A

`at()` (`schedule.Job` method), 9

C

`cancel_job()` (in module `schedule`), 7

`cancel_job()` (`schedule.Scheduler` method), 8

`CancelJob`, 8

`clear()` (in module `schedule`), 7

`clear()` (`schedule.Scheduler` method), 8

D

`day` (`schedule.Job` attribute), 9

`days` (`schedule.Job` attribute), 9

`default_scheduler` (in module `schedule`), 7

`do()` (`schedule.Job` method), 10

E

`every()` (in module `schedule`), 7

`every()` (`schedule.Scheduler` method), 8

F

`friday` (`schedule.Job` attribute), 9

H

`hour` (`schedule.Job` attribute), 9

`hours` (`schedule.Job` attribute), 9

I

`idle_seconds` (`schedule.Scheduler` attribute), 8

`idle_seconds()` (in module `schedule`), 7

J

`Job` (class in `schedule`), 8

`jobs` (in module `schedule`), 7

M

`minute` (`schedule.Job` attribute), 9

`minutes` (`schedule.Job` attribute), 9

`monday` (`schedule.Job` attribute), 9

N

`next_run` (`schedule.Scheduler` attribute), 8

`next_run()` (in module `schedule`), 7

R

`run()` (`schedule.Job` method), 10

`run_all()` (in module `schedule`), 7

`run_all()` (`schedule.Scheduler` method), 8

`run_pending()` (in module `schedule`), 7

`run_pending()` (`schedule.Scheduler` method), 8

S

`saturday` (`schedule.Job` attribute), 9

`schedule` (module), 7

`Scheduler` (class in `schedule`), 8

`second` (`schedule.Job` attribute), 9

`seconds` (`schedule.Job` attribute), 9

`should_run` (`schedule.Job` attribute), 10

`sunday` (`schedule.Job` attribute), 9

T

`tag()` (`schedule.Job` method), 9

`thursday` (`schedule.Job` attribute), 9

`to()` (`schedule.Job` method), 9

`tuesday` (`schedule.Job` attribute), 9

W

`wednesday` (`schedule.Job` attribute), 9

`week` (`schedule.Job` attribute), 9

`weeks` (`schedule.Job` attribute), 9